

# Project Survey, fundamental distributed algorithms Inria 2017

Stewart Grant

## 1 Specification and verification of Dynamic Properties in Distributed Systems

*Ozalp Babaoglu and Micheal Raynal ISIRA rennes*

### 1.1 Notes

interesting point early on that any number of events may be "irrelevant with respect to a given property" so their model simply ignores all that are not relevant. In the future consider providing a field or annotation in a messaging block, (such as govector) which would allow the isolation of messages which pertained to a given property. This would allow for a higher degree of fidelity. In addition it may be possible to associate messages with a given functionality based on their structure.

**run** a total order of events in a distributed system. This is what really happens, but is unobservable due to clock skew, and an absent centralized clock.

Examples of predicates which cannot be satisfied by a simple boolean predicate  $A \vee B \wedge C$  where  $A, B, C$  are variables.

- Load in the network is balanced.
- Resource allocation violates the at most k-out-of-n concurrent access requirement
- Message delay along the route for A to B is less than 5ms
- No more than 100 total users logged in to machines A, B, and C
- In the computation of Fig 1  $x \geq 2y$ .

**sequence spec is an ordered array of simple predicates**  
 $SS ::= SP | SS; SP$

Interval controlled sequences are a key component to realizing temporal predicates essentially the process is as follows. An observation is a set of states  $\Sigma$  on which a property can hold. As mentioned in the prior section  $SS ::= SP | SS; SP$  is a sequence of simple properties which hold over subsets of an observation. Temporal properties can be constructed by linking predicates. In

the form  $\phi; [\bar{\Theta}]; \phi$  where all such  $\phi$  are observations on which a predicate holds.

$\phi_1; [\bar{\Theta}]; \phi$  is satisfied by an observation if there exists a global state  $\Sigma^1$  in which  $\phi_1$  holds, a later global state  $\Sigma^2$  in which  $\phi_2$  holds, and  $\bar{\Theta}$  does not hold in the global states in between. Denoted in the form  $\phi_1; [\bar{\Theta}]\phi_2$ .

This paper makes the contribution of formalizing temporal strings of if *interval-constrained sequence* with the syntax  $[\bar{\Theta}_1]\phi_1; \dots; [\bar{\Theta}_m]\phi_m; [\bar{\Theta}_{m+1}]$

Key takeaway from the formulation. The difference between SS and ICS is that between states  $i_n, i_{n+1}$  there exists a  $[\bar{\theta}_n]$  which must not hold on that interval. In that way the expressiveness of ICS is greater than that of SS.

Two algorithms are proposed for solving SS **Pos** and **Def**. To compute **Def** all paths through an observation are taken and must show that the SS formula holds. In contrast SS **Pos** is verified simply by finding a single example in an observation on which the formula holds.

### 1.2 Take Away

This paper presented a novel formulation for describing temporal properties in distributed executions. It is very similar to Dinv in that it uses a lattice to formulate an execution. However, where Dinv prunes the search space by examining only strongly consistent cuts, this approach dives deeper, and actually traverses all possible interleavings in order to satisfy a predicate.

If Dinv were to take this approach it would require that a user specify their invariants ahead of time, if not then all possible observations of a run would need to be fed into Daikon.

Of specific interest is their formulation of  $[\bar{\theta}]\phi$  in that it requires some invariant to not hold before a known property in a given state. Using Dinv's ground states this same technique could still be applied, although it would be equally restricted by ground states. However, this could still be a potential extension.

The clear problem with mining invariants using this technique is that it is a search, therefore any attempt to infer invariants would result in inferring on the entire search space.

This approach does not take into account failures, it

could also be interesting to extend the expressivness of the formula to accept failures and restarts. The hardest part in this would be modifying  $\Sigma$  to accept "Global" predicates on a subset of nodes. Further it would be interesting to see if predicates could be defined which allows for the recovery of a node, by a new one introduced to the system.

Overall this is just a scheme for temporal logic detection in a system instrumented with vector clocks. It could be used for that golden key value spec  $\forall get(x) \leftarrow put(x) get(x) == get(x)$  untill  $put(x')$ . This would have been a cool path to go down if it had not allready been done.

### 1.3 Questions

- Is there a way to divide temporal search paths for mining?
- Is there a way to identify identical search paths and merge them for compression?
- Is there an obvious or non obvious way to trim a search path, or limit the lenght of a search for temporal invariant mining?
- Can Search Paths be presented to daikon in a way which preserves temporal logic?

## 2 On the Fly Testing of Regular Patterns in Distributed Computations

Eddy Fromentin, Michel Raynal Vijay k Garg, Alex Tomlinson

### 2.1 Notes

This work aim to detect regular patterns in the exectuion of distrubted systems. The first technique to note is that they reduced a distributed execution down to only events corresponding to msg receives and sends. This is the same intuition I had when initially building GoVector, and instrumenting systems for Dinv but it is done formally.

Their technique relies on labling events with a symbol in an execution language. Series of events form automata which are describable by regualr expressions. Not sure where this paper is going, but I assume that they are going to describe a partially orded regex syntax.

This paper did not go much further, the algorithm for performing on the fly regular expressions is relatively simple. Essentially every process maintains an array of the regular expression. The expression nessisarily begins at the beginning of the computation, with omits much of the

complexity incurred in the *specification and verification of dynamic properties in distributed computation* paper. Each process maintains a boolean array of the properties in the regex. When a new event occurs locally the process checks if it violates the regex. Otherwise two separete cases occur. 1) a sent message is received or a local event happens, merge the two and check if they violate any of the regex conditions. On a send event just treat it as a local event. In addition to vector clocks the array of regular experssion must be added to messages.

### 2.2 Observations

I think that regular expressions on states are interesting, but ultimatly not going to play into my work. These models could be usefull for specifiing something like TCP where the execution is known to be finite, or should be. In typical system cases though I believe that specifications need to be more expressive than regular expersions. For instance Coq and TLA+ the key languages in specification are fully fledged programming languages.

What this work does contribute is an online way to process regex. I think that this is usefull for a couple of reasons 1) I think that it could be key in generalizing some well known protocols such as leader election, where the set of known messages would fit in the vocabulary of events.

I should do a bit of follow up on this work to see if anyone built a tool which automatically translates regex into protocols rather than check a protocol with regex, but I don't think that it would yeild interesting work.

## 3 Shared Global states in Distrubted Computations

Eddy Fromentin, Michel Raynal

### 3.1 Notes

The paper begins with the same notation used in prior papers defining partial order and such. The key part of the paper is *shared global states*, which are an element in the lattice of global states such that

Let  $\Sigma = (s_1, \dots, s_i, \dots, s_n)$  be a global state  
 $\Sigma$  is shared  $\iff$

$\forall (i, j) :: (\text{prev}(s_i) \rightarrow \text{next}(s_j) \text{ or } s_j = s_i^{\text{last}} \text{ or } s_i = s_j^0)$

Essentially a shared global state is a state which has the entire global history as its transitive predissessor, or at least all of the history it could known about transitively.

The algorithm near the end of the paper shows how shared global state can be detected on the fly, however the algorithm is centralized with diminishes how interesting

it is. That stated, shared global state is near conceptually with common knowledge, in that it is the nth degree of everyone knows everything, it is the case that someone knows everything [2].

### 3.2 Observations

While it is kind of cool that the idea of *Shared Global States* has been formalized here, there is a big want with regard to purpose. Specifically when would you want to use shared global state. I can only think of Chain Replication. [3]. While shared global state is not a central tenant of the system it would achieve it quite often during execution.

With regard to Dinv this is almost certainly not the path to go down, SCC are already too much of a restriction as they do not even consider the world of *observations*. With that said GSS could be an interesting root at which to perform some sort of analysis. It serves as an obvious break point.

If I wanted to infer arbitrary temporal properties on a distributed execution I could for instance make users provide a template, search for observations from GSS roots, enumerate them in a key value store, overlay matching or equivalent states (there would be few, but some may match the template).

This is probably not the way to go either, the size of the output would be gigantic, just note that **GSS can be used for roots of computation** as they are as close to common knowledge as you are ever going to get (and it can be done on the fly)

### 3.3 Questions

- Is shared global state sufficient to coordinate locks? If not what information is required to allow for locking?
- How frequently do these happen in practice, for example in etcd? If the answer is kind of frequently there may be some interesting work to be done.

## 4 On-the-fly analysis of distributed computations

Eddy Fromentin, Claude Jard, Guy-Vincent Jourdan, Michel Raynal

### 4.1 Overview

This paper proposed a method for checking properties of distributed computations during execution. The Idea is a bit subtle, and very general. Given a standard model of

a distributed systems ie  $P_0 \dots P_n$  processes and with events  $e_0 \dots e_m$  apply to each event a work or subset of an alphabet  $\Sigma$ . As control flow passes through events on different machines check the current alphabet against a specified automata or regular expression. If an event fails to correlate with the regular expression then it has not met the specification and fails. Otherwise the letters of the event are appended to the message, and upon the next event it is checked against the regex again.

There is a little bit in the beginning about simplifying a distributed execution down to an LPO. The graph (Figure 1b) does not make a lot of sense. It seems to trim out long messages, and only be concerned with the longest message path through the execution. It was not totally clear to me why this would be used, but I guess the idea is that newer messages have some sort of precedent, and can invalidate old messages??

## 5 Efficient Distributed Detection of Conjunctions of Local Predicates

Michel Hurfin, Masaaki Mizuno, Michel Raynal, Mukesh Sigal

### 5.1 Notes

The collapse of internal events, ie the job done by the key-value store in Dinv is referred to here as an interval. More precisely "the  $x$ 'th interval of process  $P_i$  denoted by  $\theta_i^x$  is a segment of the computation that begins at event  $e_i^x$  and ends at  $e_i^{x+1}$ . This definition may be useful in describing the usage of the key value store, and this citation, along with **TODO** ▶ *on the fly analysis & shared global state and others ... would be useful. Thanks Raynal* ◀

In early Dinv I defined a total ordering on cuts. Although I never ended up using it because it was too coarse grain the same dependency is defined here. Formally: A dependency relation denoted by  $\rightsquigarrow$  is defined over the set of all consistent cuts as follows. Let

$C^x = \theta_1^{x_1}, \theta_2^{x_2}, \dots, \theta_p^{x_p}$  and  $C^y = \theta_1^{y_1}, \theta_2^{y_2}, \dots, \theta_p^{y_p}$  be two consistent cuts, then:  $C^x \rightsquigarrow C^y \iff (C^x \neq C^y) \wedge (\forall k, 1 \leq k \leq p, x_k \leq y_k)$

This algorithm is very technical, but somewhat similar to Dinv. Essentially it aims to show that at some point in relative time a conjunction of local predicates held. The algorithm involves two vector clocks and a boolean vector denoting predicate satisfiability of on a process.

I think the algorithm goes like this. If a predicate is true, set  $P_i$  to true. If you get a message with  $P_j = \text{true}$  update the predicate array. Mark the current vector time, and keep track of the first point in time the predicate was denoted true. If  $\theta$  becomes false set all incoming and

outgoing predicate messages to false. The predicate becomes true if there is a causal path across all processes such that  $\forall i \theta_i = \text{true}$ . It is not super complicated, the hard part is really just that it is done on the fly.

I think that this algorithm is a bit pessimistic, using a lattice I think that these predicates could be detected more easily.

## 5.2 Observations

This work is very specific. It aims to detect global predicates online, with significant computational overhead. I think that the key takeaway from the algorithm is the technique for detecting causal intervals online. Essentially taint tracking and merging across multiple processes to show that some property  $\theta$  is satisfied globally.

This work seeks to show that a given predicate is recognized globally and does not have any interaction with a state machine. Could this online technique be used as an atom to verify sequences in computation, and allow for decentralized distributed computation? For instance would it be possible to encode a bunch of tasks, and have a single process issues a continue message once the task is complete? If the satisfaction of a given task is monotonic then I believe that the answer is yes.

## 5.3 Questions

- Could this technique be used to easily distribute SAT solving? my initial instinct is no, because there is no obvious way to do back propagation.
- Is this applied anywhere? For instance this could be used to show something like termination in Hadoop without a specified leader. With that stated it would be a bit of overkill.
- could this be used to simplify termination in graph processing? Is this a more robust solution to predicate detection than the *White/Black* graph coloring that pregal uses.
- Is there a recovery mechanism for satisfied predicates? It seems like in the worst case only a single process will learn that the predicate is satisfied.

# 6 Detecting Atomic Sequences of Predicates in Distributed Computations, 1993

Hichel Hurfin, Noel Plouzeau, Micheal Raynal

## 6.1 Notes

The goal of this paper is to detect sequences of atomic properties. I assume that these would be something similar to the detection of distributed mutual exclusion. The key insight that was presented in the abstract is that these properties are not monotonic. Perhaps this means that atomic properties may hold on the **POS** predicate? I'm not sure because for the sake of guaranteed atomicity **DEF** seems to be a necessary constraint.

Key note, the definition of "Atomicity" in the case of this paper is a sequence of events, of which none violate a specific property ie  $CA = \neg NP_1 \wedge NP_2 \wedge \dots \wedge NP_n$ . Where each  $NP_i$  denotes the property on a set of processes.

The real purpose of this paper is to count the number of causal paths which satisfy a prefix of the form  $do_0; [don't_0]do_1; [don't_1] \dots do_n$ . This differs from the specifying systems paper a little bit, essentially in the way that they are counting the number of occurrences.

## 6.2 Observations

An anti observation is, what is the purpose of this algorithm. It seems to be a generalization of the specification language, defined in Specification and Verification of dynamic properties, but it is simply a counting function. I suppose that the interesting bit comes from handling cases where predicates are invalidated due to being non-atomic however, it was unclear to me why a counter of these properties was wanted and not simply a flag.

This algorithm is extremely clean. It would be interesting to see if it could be extended to define a generalized decentralized computation framework similar to Dryad but with a decentralized compute. Essentially all computations would be defined as a set of predicates and antipredicates, for instance, compute a property on some large graph, and don't compute where any other node has. Instances of conflict could be counted and reset. Individual functions could be specified in this way, and entire programs could be written as a hierarchy of predicates, and loops defined on them. Figuring out how to work in conditionals would be a problem onto itself.

## 6.3 Questions

- What common distributed debugging tasks actually require counts of invariant violations.
- Is it reasonable to define atomic predicates, with the exception of locking?
- Can this be leveraged in reverse, and turned into a control mechanism, rather than a checking algorithm.

## 7 Some Optimal Algorithms for Decomposed Partially Ordered Sets

Vijay K. Garag

### 7.1 Notes

I'm reading this one as both an introduction to Garag and as a bit of a break from the proof heavy papers.

This paper already seems like it is going to be a good resource, it has an algorithm for detecting consistent cuts in  $n^2m$  time ... if this is really the case, it may be possible just cut the crap on dinv and really speed it up.

As usual my hopes are dashed this is an early algorithm which only detects a single consistent cut, not all of them like I had wanted. However, the online variant may be useful for making Dinv a streaming service at some point.

The second algorithm in this paper checks if a set of given antichains compose are constrained by a total ordering. The algorithm does so in  $O(mn \log n)$  time. Essentially it continuously merges sets of antichains until it finds a conflict where some set of antichains has events / elements  $s$  and  $t$  such that  $s \parallel t$ . It is a very nice algorithm, but given that Dinv never encounters total orderings I don't think that it will come in handy.

### 7.2 Observations

The online detection of consistent cuts could be used as a front end for dinv easily, but the cost of checking all consistent cuts is still exponential.

### 7.3 Questions

- What is the purpose of finding largest antichains? Potentially it shows that the largest number of chains which can be composed is  $n$  but it is still not clear where this would be applied. Is this to show the complexity during distributed debuggin?

## 8 Predicate Detection for Parallel Computation with Locking Constraints

Yen-Jung Chang, Vijay K. Garag

### 8.1 Notes

Initially this paper proposes an analysis technique which analyzes an exponential number of posets in order to check if locking predicates hold. It seems as if it is going to build posets out of code itself.

After getting through the first 5 pages I was very lost. This paper it seems is too theoretical for me.

**TODO** ▶ Return to this paper after reading 15 more theoretical papers by Garag ◀

## 9 Modeling, Analyzing and Slicing Periodic Distributed Computations

Vijay K. Garag, Anurag Agarwal, Vinit Ogale

### 9.1 Notes

This paper proposes an analysis technique for decomposing distributed executions. It proposes a very interesting model, in which DAG's are infinite. Their proposed infinite DAG is referred to as a  $d$ -diagram and essentially models looping computation post predicate on an infinite timeline. They use this model to verify temporal properties such as liveness.

Not the usage of *cut frontier*. This is proper usage and should be integrated into the Dinv paper if possible.

A very interesting addition on page 14 is the formalization of recursive vector clock. It is simple algebra, but it states simply that given some initial configuration  $C$  lets say, and a continuing recursive computation, the vector clock values for all events, of future computations can be computed a priori.

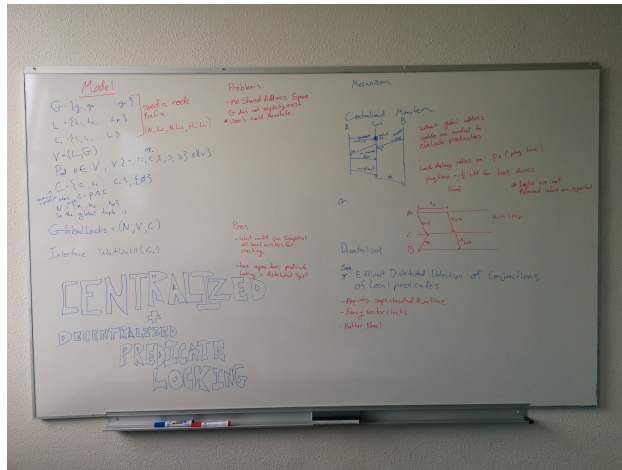
The paper continues by formulating how to verify liveness properties

### 9.2 Observations

My immediate impression of this paper is that it falls short in its model of a distributed execution. It only defines an infinite execution as a prefix followed by a periodic set of recurrent transitions. It seems that this model is strictly limited to a single periodic computation, and not a general computation which must adapt to interaction.

For instance this model could be used to analyze a key value store which only server puts, or gets, or a finite





**Figure 2:** Whiteboard of potential distributed (centralized and decentralized) predicate locking

the expectation was that locks would last a long time [1]. Has work been done where only waiting threads are issued signals based on predicate evaluation, and not just locks?

- Can the evaluation of predicate be done without a centralized source, and in an online fashion?

## References

- [1] M. Burrows. The chubby lock service for loosely-coupled distributed systems. In *Proceedings of the 7th Symposium on Operating Systems Design and Implementation*, OSDI '06, pages 335–350, Berkeley, CA, USA, 2006. USENIX Association.
- [2] J. Y. Halpern and Y. Moses. Knowledge and common knowledge in a distributed environment. *J. ACM*, 37(3):549–587, July 1990.
- [3] R. van Renesse and F. B. Schneider. Chain replication for supporting high throughput and availability. In *Proceedings of the 6th Conference on Symposium on Operating Systems Design & Implementation - Volume 6*, OSDI'04, pages 7–7, Berkeley, CA, USA, 2004. USENIX Association.